

Recurrent Neural Networks for Natural Language Processing



Data Analytics and Machine Learning

Overview



10:00 – 12:00
LECTURE



12:00 – 1:50
LAB SESSION



1:50 – 2:00
SESSION WRAP UP

Lesson Objectives

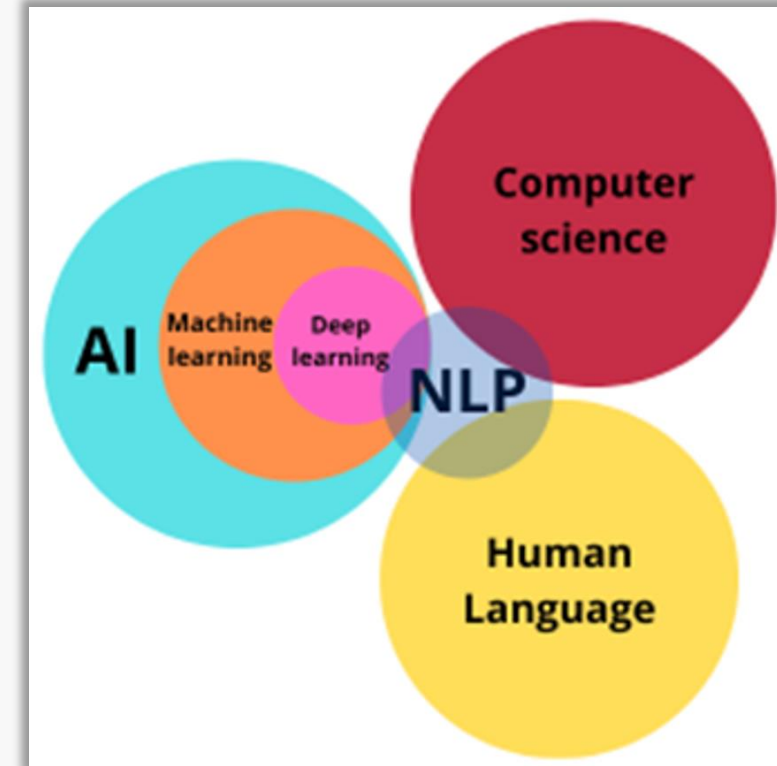
- Understand how machine learning models handle non-numeric data
- Understand which scenarios these might be applied to
- Work through a practical example

Keyword	Description
Machine Learning	A type of artificial intelligence that enables computers to learn from data and make decisions or predictions without being explicitly programmed.
Classification	This is a type of supervised learning task where the goal is to predict a categorical target variable.
Hyperparameter	These are parameters of the learning algorithm itself, not derived from the data, that need to be set before training the model.
Neural Network	A set of algorithms modelled after the human brain, designed to recognize patterns. They interpret sensory data through a kind of machine perception, labelling or clustering raw input.
Recurrent Neural Networks (RNN)	A type of artificial neural network designed to process sequential data. RNNs take information from prior inputs to influence the current input and output.
Long Short-Term Memory (LSTM)	A subtype of RNN that is designed to remember past information while forgetting irrelevant parts, making it effective for tasks involving sequential data.
Gated Recurrent Unit (GRU)	A type of RNN that operates similarly to LSTM but has fewer parameters, making it computationally more efficient.
Text Classification	A machine learning technique that assigns predefined categories to open-ended text, helping to organize and structure any kind of text data.
Early Stopping	A form of regularization used in machine learning to prevent overfitting. It allows training to stop once the model's performance stops improving on a hold-out validation dataset.
Optimizer	A function or algorithm that adjusts the attributes of a neural network, such as weights and learning rates, to minimize the loss function and improve the model's performance.

Introduction to Natural Language Processing

NLP - Natural Language Processing

- A branch of computer science
- Integration of computing and human language
 - Computational linguistics
 - Statistical Computer Science
 - Machine learning
 - Deep learning models
- Facilitates understanding of full meaning including intent and sentiment of speaker or writer



Natural Language Processing (NLP)

- Translate text
- Summarise large volumes of text
- Respond to spoken commands
 - Voice commanded GPS
 - Digital Assistants
 - Speech-to-text dictation
 - Customer service chatbots
- Increasingly being used in enterprise solutions
 - Streamline business operations
 - Simplify mission-critical processes



NLP – Human Language difficulties

- Human language is filled with **ambiguities**
 - Hard to create software that accurately determines:
 - Intended meaning
 - Context
 - Emotion
 - **Irregularities** take humans years to learn
 - Must be taught to natural language-driven applications from the start for them to be useful
- **Challenges** include
 - Homonyms
 - Homophones
 - Sarcasm
 - Idiom
 - Metaphors
 - Grammar
 - Usage exceptions
 - Used in a different manner to their definition
 - Sentence structure variations e.g. to avoid monotony or to provide emphasis

NLP - Tasks

Restricted - Other

Speech recognition

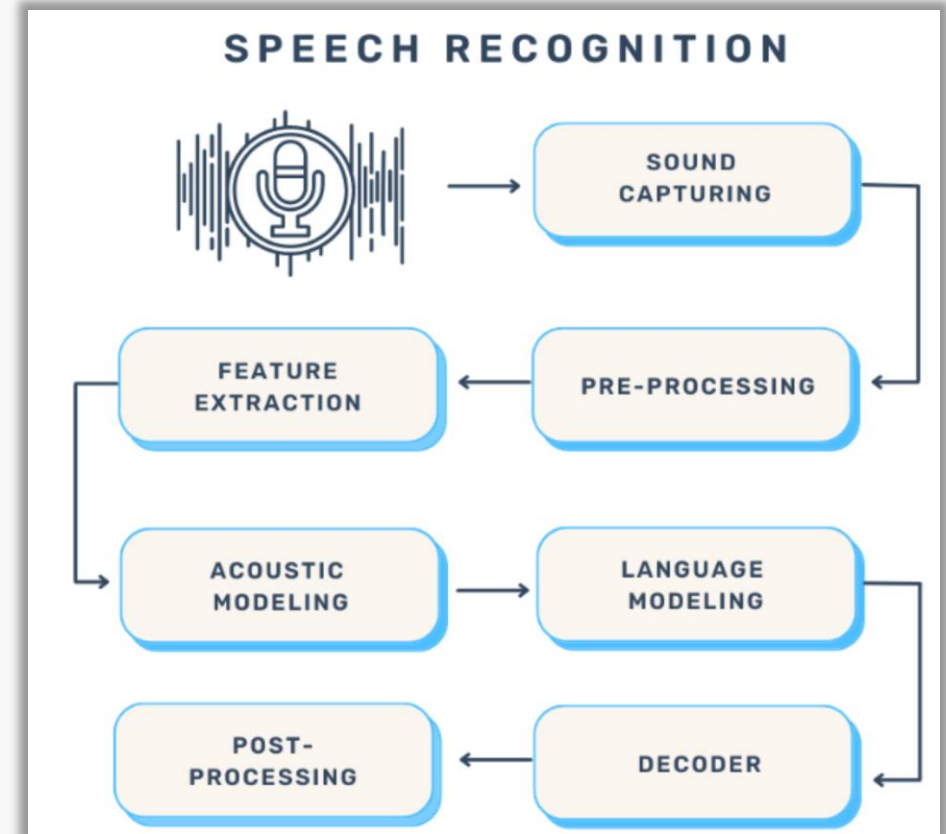
- Converts voice data into text data
 - **Implementation**
 - Machine learning algorithms to interpret human speech
 - **Example**
 - Voice-operated GPS systems

Part of speech tagging

- Determines speech of word based on its use and context
 - **Implementation**
 - Uses linguistic rules and statistical models
 - **Example**
 - Identifies 'make' as a verb 'I can make a paper plane'

Word sense disambiguation

- Selects the meaning of a word based on context
 - **Implementation**
 - Uses semantic analysis algorithms
 - **Example**
 - Distinguishes the meaning of 'make' in 'make the grade'
 - vs. 'make a bet'



NLP - Tasks

Named entity recognition (NEM)

- Identifies words or phrases
 - As useful entities
 - **Implementation**
 - Machine learning models
 - Trained on annotated data
 - **Example**
 - Identifies 'Kentucky' as a location
 - Or 'Fred' as a man's name

Co-reference resolution

- Identifies if two words refer to the same entity
 - **Implementation**
 - Uses rule-based methods
 - And machine learning models
 - **Example**
 - Determines that 'she' refers to 'Mary'
 - In a given context



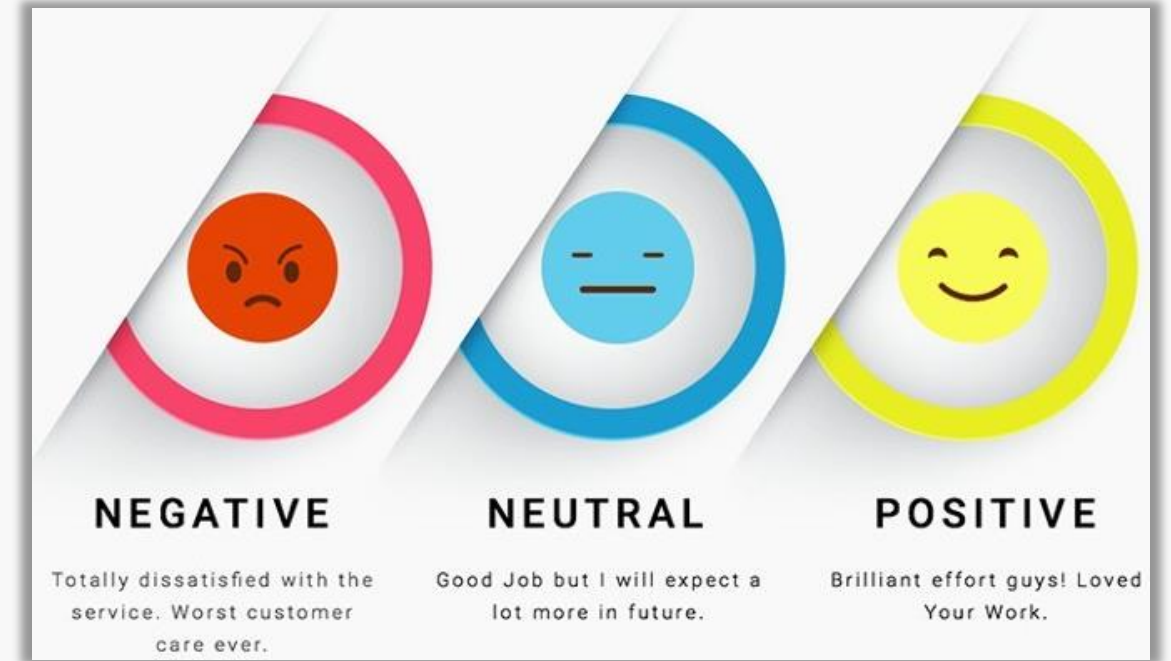
NLP - Tasks

Sentiment analysis

- Extracts subjective qualities like attitudes and emotions from text
 - **Implementation**
 - NLP Text analysis
 - Computational linguistics
 - **Example**
 - Determines user sentiment
 - From product reviews

Natural language generation

- Converts structured information
 - Into human language
 - **Implementation**
 - Uses
 - Templates
 - Rules
 - Machine learning models
 - **Example**
 - A weather app that generates a weather report
 - From meteorological data



NLP – Use Cases

Social media sentiment analysis

- Extracts attitudes and emotions
 - from social media
 - Posts
 - Responses
 - Reviews, etc
- Implementation
 - Natural language processing
 - Text analysis
 - Computational linguistics
- Example
 - Analysing customer sentiment
 - Towards products or promotions on social media platforms

Text summarisation

- Creates summaries
 - Of large volumes of digital text
- Implementation
 - Semantic reasoning
 - Natural language generation
 - To add to summaries
 - Useful context
 - Conclusions
- Example
 - Summarising news articles or research papers
 - For quick reading

NLP – Good Libraries

Hugging Face Transformers

- Pre-trained models for tasks on text, vision, audio

spaCy

- Supports tokenization and training for 60+ languages

Fairseq

- Train custom models
- Translation, summarisation, language modelling, etc

Jina

- Building scalable neural search applications

Gensim

- Topic modelling, document indexing, and similarity retrieval with large corpora

NLTK (Natural Language Toolkit)

- Platform for building Python programs to work with human language data

TextBlob

- A simple API for common NLP tasks
- Part-of-speech tagging, noun phrase extraction, sentiment analysis, classification, translation, etc.

CoreNLP

- Group of NLP Programs
- Tokenization, part-of-speech tagging, lemmatization, etc

Polyglot

- Perform different NLP operations

Scikit-learn

- Intuitive class methods and numerous algorithms
- To build machine learning models

Pattern

- Implementing Natural Language processing tasks
- Text Mining, NLP, and Machine Learning

NLP Sentiment Analysis and Classification with TensorFlow

Data Pre-processing

- **Download the dataset**
 - From Kaggle
 - <https://www.kaggle.com/datasets/lakshmi25npathi/imdb-dataset-of-50k-movie-reviews>
 - Place it in the project's home directory
- **Read the dataset**
 - Convert the sentiment column
 - To numeric values
 - For binary classification
 - Use np.where()
 - 'positive' sentiment = 1
 - 'negative' sentiment = 0
- **Convert the labels and reviews to NumPy arrays**
 - Pre-processing methods favour arrays
 - Over Pandas series

Restricted - Other

Use pandas to load the IMDB review dataset

```
1 reviews = pd.read_csv('IMDB Dataset.csv')
2 reviews.head()
```

✓ 0.5s

	review	sentiment
0	One of the other reviewers has mentioned that ...	positive
1	A wonderful little production. The...	positive
2	I thought this was a wonderful way to spend ti...	positive
3	Basically there's a family where a little boy ...	negative
4	Petter Mattei's "Love in the Time of Money" is...	positive

Convert the sentiment column to numeric values for binary classification

Regard 'positive' sentiment as 1 and 'negative' sentiment as 0 using np.where()

```
1 # Regard 'positive' sentiment as 1 and 'negative' sentiment as 0
2 reviews['sentiment'] = np.where(reviews['sentiment'] == 'positive', 1, 0)
```

✓ 0.0s

Convert the labels and reviews(sentences) to NumPy arrays

Note: Pre-processing methods favor arrays over pandas series

```
1 # Convert the labels and reviews(sentences) to NumPy arrays
2 sentences = reviews['review'].to_numpy()
3 labels = reviews['sentiment'].to_numpy()
```

✓ 0.0s

Data Pre-processing

- **Train/Test Split**
- Split the dataset (train/test split) before any pre-processing
- 75:25 split
- Dataset is 50,000 reviews
 - Training model using 37500 reviews
 - Testing model accuracy using the unseen 12500 reviews

Split the dataset into training and test instances before any pre-processing

Use a 75:25 split for training and testing data, respectively

Dataset is 50,000 reviews

Training LSTM model using 37500 reviews

Testing model accuracy using the unseen 12500 reviews

```
1 X_train, X_test, y_train, y_test = train_test_split(sentences, labels, test_size=0.25)
2 print("Training Data Input Shape: ", X_train.shape)
3 print("Training Data Output Shape: ", y_train.shape)
4 print("Testing Data Input Shape: ", X_test.shape)
5 print("Testing Data Output Shape: ", y_test.shape)
```

✓ 0.0s

```
Training Data Input Shape: (37500,)
Training Data Output Shape: (37500,)
Testing Data Input Shape: (12500,)
Testing Data Output Shape: (12500,)
```


Data Pre-processing

Tokenisation on the entire text corpus

- Includes all the training data reviews

Convert textual data

- Reviews
 - Into numeric values
 - To build a mathematical model
- Specify **vocabulary size**
 - Tokenisation of training data
- Consider the **first 10000 words**
 - Based on frequency
 - In the training data
- Specify **oov_tok**
 - As <OOV>
 - Replaces any unknown word in the text corpus

```
1 # Set the vocabulary size to 10000
2 vocab_size = 10000
3 # Set the out-of-vocabulary token to "<OOV>"
4 oov_tok = "<OOV>"
5 # Initialise the tokenizer with the specified vocabulary size and OOV token
6 tokenizer = Tokenizer(num_words=vocab_size, oov_token=oov_tok)
```

Data Pre-processing

Tokenise sentences

- Into a set of **individual words** during tokenisation
- Calculate statistical features for each word
- **word_counts**
 - Dictionary of words with word count in the entire text corpus
- **word_docs**
 - Dictionary of words depicting the number of documents in the text corpus containing a specific word
- **word_index**
 - A unique index assigned to a dictionary of words
- **document_count**
 - Number of documents used for fitting the tokenizer

```
1 # Fit the tokenizer on the training data
2 tokenizer.fit_on_texts(X_train)
3
4 # Print the number of documents used for fitting the tokenizer
5 print("Number of Documents: ", tokenizer.document_count)
6
7 # Print the number of words in the tokenizer
8 print("Number of Words: ", tokenizer.num_words)
```

✓ 4.9s

```
Number of Documents:  37500
Number of Words:  10000
```

Data Pre-processing

Hyperparameters for tokeniser

- Fit the hyperparameters for Tokenizer() on the training data
 - fit_on_texts()
- Visualise:
 - the count of each word in the overall dictionary
 - The number of documents containing a specific word
- Convert each textual review into a numerical sequence using the fitted tokenizer

```
1 # Fit the tokenizer on the training data
2 tokenizer.fit_on_texts(X_train)
3
4 # Print the number of documents used for fitting the tokenizer
5 print("Number of Documents: ", tokenizer.document_count)
6
7 # Print the number of words in the tokenizer
8 print("Number of Words: ", tokenizer.num_words)
```

✓ 4.9s

Number of Documents: 37500
Number of Words: 10000

```
1 # Print the word counts in the tokenizer
2 tokenizer.word_counts
```

✓ 0.0s

```
OrderedDict([('once', 3425),
             ('upon', 1354),
             ('a', 242241),
             ('time', 18773),
             ('there', 23618),
             ('was', 71739),
             ('science', 790),
             ('fiction', 727),
             ('author', 331),
             ('named', 1151),
             ('h', 344),
             ('beam', 25),
             ('piper', 108),
             ('who', 30325),
             ('wrote', 821),
             ('classic', 2652),
             ('book', 3474),
             ('little', 9190),
             ('fuzzy', 83),
             ('which', 17549),
             ('about', 25496),
             ('man', 8248),
             ('discovering', 129),
             ('race', 589),
             ('of', 216196),
             ...
             ('heels', 87),
             ('private', 397),
             ('detective', 626),
             ('yarn', 48),
             ...])
```

```
1 # Print the word documents in the tokenizer
2 tokenizer.word_docs
```

✓ 0.0s

```
defaultdict(int,
              {'where': 7088,
               'man': 5707,
               'about': 15482,
               'who's': 895,
               'ewoks': 21,
               'adorable': 137,
               'for': 26690,
               'fuzzy': 78,
               'is': 33473,
               'race': 439,
               'to': 35195,
               'org': 17,
               'this': 33988,
               'today's': 360,
               'free': 946,
               'blatant': 156,
               'science': 600,
               'before': 5347,
               'died': 712,
               'mr': 1424,
               'and': 36218,
               'project': 669,
               'upon': 1232,
               'take': 4453,
               ...
               'priceless': 134,
               'mark': 793,
               'strongest': 106,
               'producing': 164,
               ...})
```

Data Pre-processing

Convert training data reviews

- Convert each review in the training data
 - Into a numerical sequence
 - For further training purposes
- Note
 - Each review has different lengths of words
 - Will produce diverse numeric sequence lengths

```
1 # Convert the tokenised training data into sequences
2 train_sequences = tokenizer.texts_to_sequences(X_train)
3
4 # Print the first sequence
5 print(train_sequences[0])
```

✓ 3.1s

[281, 684, 4, 56, 47, 14, 4, 1092, 1166, 2266, 792, 2194, 1, 5385, 37, 1060, 4, 361, 278, 792,

Data Pre-processing

Limit Sequence Lengths

- To a constant value for each review
- Set a nominal **sequence length**
 - 200 for each review
- **Truncate** numerical sequences
 - Lengths greater than 200
- **Pad** sequences
 - Lengths smaller than 200
 - With zeros
- Set the **sequence padding** for numerical sequences
 - Of textual reviews
- **Repeat** the same pre-processing steps
 - For **test data**
 - after training data complete
- **Complete pre-processing of the textual reviews**
 - Tokenisation, sequence conversion, and padding

```

1 # Convert the tokenised training data into sequences
2 train_sequences = tokenizer.texts_to_sequences(X_train)
3
4 # Print the first sequence
5 print(train_sequences[0])
✓ 3.1s

```

[281, 684, 4, 56, 47, 14, 4, 1092, 1166, 2266, 792, 2194, 1, 5385, 37, 1060, 4, 361, 278, 792, 120, 6531, 61, 14,

```

1 sequence_length = 200
2 train_padded = pad_sequences(train_sequences, maxlen=sequence_length, padding='post', truncating='post')
✓ 0.3s

```

```

1 # Convert the tokenised test data into sequences
2 test_sequences = tokenizer.texts_to_sequences(X_test)
3
4 # Pad the sequences to ensure uniform length, truncating longer ones and padding shorter ones with zeros
5 test_padded = pad_sequences(test_sequences, maxlen=sequence_length, padding='post', truncating='post')

```

Building an RNN using TensorFlow

Initiate a new Sequential model

- This model serves as a linear stack of layers
 - In the neural network
- Each layer's output
 - Is the input for the next layer
 - Last layer outputs
 - Prediction label
- Use this model to embed the layers of the LSTM (Long Short-Term Memory) network
- The LSTM layers can be added to this model

```
1 # Convert the tokenised training data into sequences
2 train_sequences = tokenizer.texts_to_sequences(X_train)
3
4 # Print the first sequence
5 print(train_sequences[0])
```

✓ 3.1s

[281, 684, 4, 56, 47, 14, 4, 1092, 1166, 2266, 792, 2194, 1, 5385, 37, 1060, 4, 361, 278, 792,

Building an RNN using TensorFlow

Add an embedding layer to the model

- This layer converts each word
 - Into a dense vector
 - Of embedding dimensions
 - Hyperparameters of the layer
- Set **vocabulary size** and **sequence length**
 - For each review
- Add a **Bidirectional()** layer
- Add a **LSTM layer to the model**
 - Set a unit size in the LSTM layer

Note: Bidirectional LSTM remembers output from:

- Past to future
- And from future to past
- More robust models for time series analysis

```

1 # Set the embedding dimension to 16
2 embedding_dim = 16
3 # Add an Embedding layer to the model
4 model.add(Embedding(vocab_size, embedding_dim, input_length=sequence_length))
5
6 # Set the LSTM output to 32
7 lstm_out = 32
8 # Add a Bidirectional LSTM layer to the model
9 model.add(Bidirectional(LSTM(lstm_out)))
10 # Add two Dense layers to the model with 'relu' and 'sigmoid' activation functions respectively
11 model.add(Dense(10, activation='relu'))
12 model.add(Dense(1, activation='sigmoid'))
13 # Compile the model with binary crossentropy loss function, adam optimizer, and accuracy metrics
14 model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
15 # Print a summary of the model
16 model.summary()
17
18

```

✓ 0.6s

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
embedding (Embedding)	(None, 200, 16)	160000
bidirectional (Bidirectional)	(None, 64)	12544
dense (Dense)	(None, 10)	650
dense_1 (Dense)	(None, 1)	11
=====		

Total params: 173205 (676.58 KB)

Trainable params: 173205 (676.58 KB)

Non-trainable params: 0 (0.00 Byte)

Building an RNN using TensorFlow

- **Add two Dense layers to the model**
 - Specify activation functions
- **Add a fully connected layer**
 - 10 units
 - 'relu' activation
- **Add an output layer**
 - 1 unit
 - 'sigmoid' activation
- **output layer**
 - Outputs probability input belongs to
 - 1 (positive) using the sigmoid filter

```

1 # Set the embedding dimension to 16
2 embedding_dim = 16
3 # Add an Embedding layer to the model
4 model.add(Embedding(vocab_size, embedding_dim, input_length=sequence_length))
5
6 # Set the LSTM output to 32
7 lstm_out = 32
8 # Add a Bidirectional LSTM layer to the model
9 model.add(Bidirectional(LSTM(lstm_out)))
10 # Add two Dense layers to the model with 'relu' and 'sigmoid' activation functions respectively
11 model.add(Dense(10, activation='relu'))
12 model.add(Dense(1, activation='sigmoid'))
13 # Compile the model with binary crossentropy loss function, adam optimizer, and accuracy metrics
14 model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
15 # Print a summary of the model
16 model.summary()
17
18

```

✓ 0.6s

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
embedding (Embedding)	(None, 200, 16)	160000
bidirectional (Bidirectional)	(None, 64)	12544
dense (Dense)	(None, 10)	650
dense_1 (Dense)	(None, 1)	11
=====		

Total params: 173205 (676.58 KB)
 Trainable params: 173205 (676.58 KB)
 Non-trainable params: 0 (0.00 Byte)

Building an RNN using TensorFlow

- **Compile the model**
 - Optimises the binary_crossentropy
 - During training
- **'adam' optimiser**
 - Minimises loss value
 - By tweaking the weights
 - During the training phase
 - Tries to find the global minima
 - For the loss value
 - Across all the local minima
- **'accuracy' of the model**
 - Reported for each training batch/epoch
 - Gauge the convergence
 - Of the neural network
- **Visualise** the summary of the LSTM model

```

1 # Set the embedding dimension to 16
2 embedding_dim = 16
3 # Add an Embedding layer to the model
4 model.add(Embedding(vocab_size, embedding_dim, input_length=sequence_length))
5
6 # Set the LSTM output to 32
7 lstm_out = 32
8 # Add a Bidirectional LSTM layer to the model
9 model.add(Bidirectional(LSTM(lstm_out)))
10 # Add two Dense layers to the model with 'relu' and 'sigmoid' activation functions respectively
11 model.add(Dense(10, activation='relu'))
12 model.add(Dense(1, activation='sigmoid'))
13 # Compile the model with binary crossentropy loss function, adam optimizer, and accuracy metrics
14 model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
15 # Print a summary of the model
16 model.summary()

```

✓ 0.6s

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
embedding (Embedding)	(None, 200, 16)	160000
bidirectional (Bidirectional)	(None, 64)	12544
dense (Dense)	(None, 10)	650
dense_1 (Dense)	(None, 1)	11
=====		

Total params: 173205 (676.58 KB)
 Trainable params: 173205 (676.58 KB)
 Non-trainable params: 0 (0.00 Byte)

Building an RNN using TensorFlow – EarlyStopping()

- **EarlyStopping()**
 - Halts model training
 - After the model fails to minimise
 - The validation loss value
 - After a set number of epochs
- Helps **avoid overfitting** the model
 - On the training data
- **ModelCheckpoint()**
 - Monitor the loss after each epoch
 - Save the best model
 - In terms of validation loss

```
1 # Set the checkpoint file path to the current working directory
2 checkpoint_filepath = os.getcwd()
3
4 # Create a ModelCheckpoint callback that saves the best model in terms of validation loss
5 model_checkpoint_callback = ModelCheckpoint(filepath=checkpoint_filepath,
6                                             save_weights_only=False,
7                                             monitor='val_loss',
8                                             mode='min',
9                                             save_best_only=True)
10
11 # Create a list of callbacks including EarlyStopping and ModelCheckpoint
12 callbacks = [EarlyStopping(patience=2), model_checkpoint_callback]
```

Building an RNN using TensorFlow – Fit Model

- **Fit the model**
 - Set the number of epochs
 - Network trained this amount
 - Number of epochs needed
 - Often unknown
 - Requires an educated guess
 - And tweaking
 - Maximum of 10 epochs
 - In example
- **Set the validation data**
 - Monitor the loss on the validation dataset
- **Halt the training**
 - If the validation loss is not minimised
 - For two consecutive epochs
 - Specified in the callback
 - Model training may halt before reaching 10 epochs
 - If the validation loss does not improve

```

1 # Fit the model on the training data for a maximum of 10 epochs
2 # monitoring the loss on the validation dataset
3 history = model.fit(train_padded,
4                       y_train, epochs=10,
5                       validation_data=(test_padded,
6                                       y_test),
7                       callbacks=callbacks)

```

Building an RNN using TensorFlow – Model Accuracy

- Model training halted
 - After 5 epochs
 - Loss did not improve
 - After Epoch 3
- Model parameters
 - Saved in 'history' variable
- Achieved 86% validation accuracy
 - On the IMDB review dataset
 - By training a simple bidirectional LSTM network
- Accuracy could be improved
 - Using back-to-back LSTM layers
 - Or using increased word dictionary

```

1 history = model.fit(train_padded,
2                     y_train,
3                     epochs=10,
4                     validation_data=(test_padded,
5                                     y_test),
6                     callbacks=callbacks)
✓ 6m 44.8s

```

```

Epoch 1/10
1172/1172 [=====] - 91s 74ms/step - loss: 0.6494 - accuracy: 0.6212 - val_loss: 0.5022 - val_accuracy: 0.7782
Epoch 2/10
1172/1172 [=====] - 85s 72ms/step - loss: 0.4619 - accuracy: 0.7881 - val_loss: 0.3759 - val_accuracy: 0.8459
Epoch 3/10
1172/1172 [=====] - 95s 81ms/step - loss: 0.3061 - accuracy: 0.8778 - val_loss: 0.3214 - val_accuracy: 0.8628
Epoch 4/10
1172/1172 [=====] - 67s 57ms/step - loss: 0.2497 - accuracy: 0.9039 - val_loss: 0.3566 - val_accuracy: 0.8443
Epoch 5/10
1172/1172 [=====] - 67s 57ms/step - loss: 0.2222 - accuracy: 0.9162 - val_loss: 0.3429 - val_accuracy: 0.8664

```

```

1 metrics_df = pd.DataFrame(history.history)
2 print(metrics_df)
✓ 0.0s

```


	loss	accuracy	val_loss	val_accuracy
0	0.649434	0.621227	0.502158	0.77824
1	0.461919	0.788107	0.375894	0.84592
2	0.306105	0.877813	0.321435	0.86280
3	0.249678	0.903947	0.356642	0.84432
4	0.222168	0.916160	0.342907	0.86640

Building an RNN using TensorFlow – Model Loss

- Visualise the
- Training/testing data
 - Over the number of epochs
 - Using matplotlib

```
1 plt.figure(figsize=(10,5))
2 plt.plot(metrics_df.index, metrics_df.loss)
3 plt.plot(metrics_df.index, metrics_df.val_loss)
4 plt.title('IMBD Reviews Sentiment Analysis Model Loss over Epochs')
5 plt.xlabel('Epochs')
6 plt.ylabel('Binary Crossentropy')
7 plt.legend(['Training Loss', 'Validation Loss'])
8 plt.show()
```

✓ 0.2s




Building an RNN using TensorFlow – Model Accuracy

- Visualise the accuracy
- Training/testing data
 - Over the number of epochs
 - Using matplotlib

```
1 plt.figure(figsize=(10,5))
2 plt.plot(metrics_df.index, metrics_df.accuracy)
3 plt.plot(metrics_df.index, metrics_df.val_accuracy)
4 plt.title('IMBD Reviews Sentiment Analysis Model Accuracy over Epochs')
5 plt.xlabel('Epochs')
6 plt.ylabel('Accuracy')
7 plt.legend(['Training Accuracy', 'Validation Accuracy'])
8 plt.show()
```

✓ 0.1s



Session Review

- Understand how machine learning models handle non-numeric data
- Understand which scenarios these might be applied to
- Work through a practical example